

Package Performance

Tuning Packages for Optimal Performance

Stacia Misner

blog.datainspirations.com

smisner@datainspirations.com

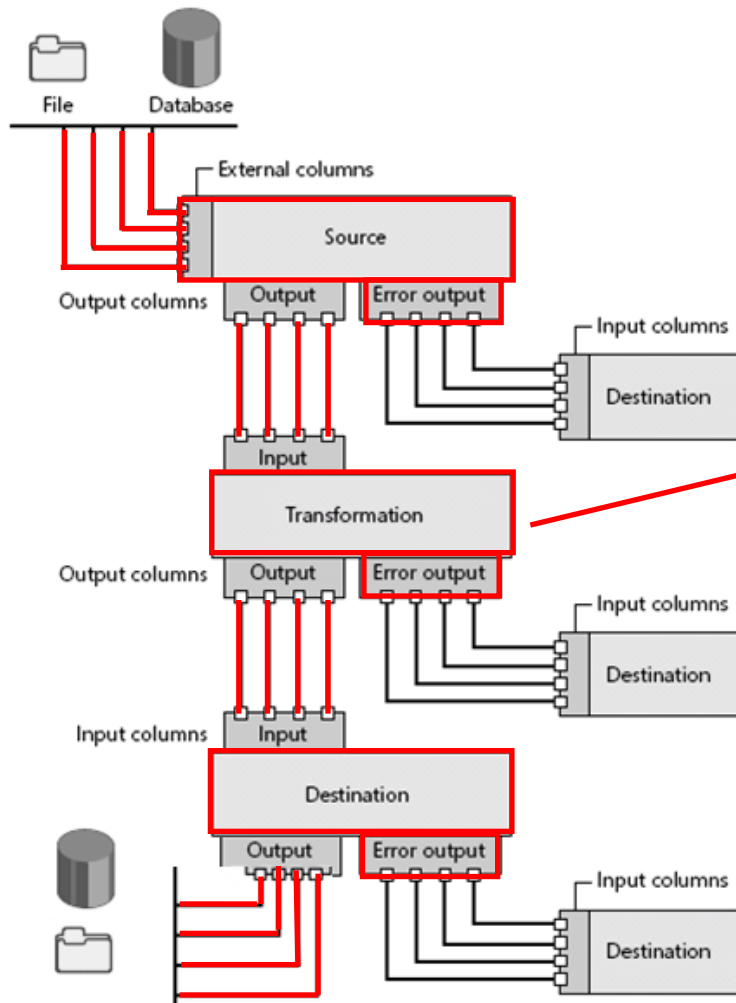


pluralsight 
hardcore developer training

Overview

- Pipeline Buffer Architecture
- Execution Trees and Execution Threads
- Package Performance Optimization
- Performance Monitoring

Pipeline Buffer Architecture



Buffer

	ProductID	Name	ProductNumber	MakeFlag	FinishedGoodsFlag	Color	SafetyStockLevel	ReorderPoint	StandardCost	ListPrice	ce
1	1	Adjustable Race	AR-5381	0	0	NULL	1000	750	0.00	0.00	
2	2	Bearing Ball	BA-8327	0	0	NULL	1000	750	0.00	0.00	
3	3	BB Ball Bearing	BE-2349	1	0	NULL	800	600	0.00	0.00	
4	4	Headret Ball Bearings	BE-2908	0	0	NULL	800	600	0.00	0.00	
5	316	Blade	BL-2036	1	0	NULL	800	600	0.00	0.00	
6	317	LL Crankarm	CA-5965	0	0	Black	500	375	0.00	0.00	
7	318	ML Crankarm	CA-6738	0	0	Black	500	375	0.00	0.00	
8	319	HL Crankarm	CA-7457	0	0	Black	500	375	0.00	0.00	

Buffers

Location of BLOB data types
storage in buffer not BLOB

DT_I4 DT_WSTR DT_WSTR DT_BOOL DT_BOOL DT_WSTR DT_I2 DT_I2 DT_CY DT_CY
4 50 25 15 2 2 8 8

	ProductID	Name	ProductNumber	MakeFlag	FinishedGoodsFlag	Color	SafetyStockLevel	ReorderPoint	StandardCost	ListPrice
1	1	Adjustable Race	AR-5381	0	0	NULL	1000	750	0.00	0.00
2	2	Bearing Ball	BA-8327	0	0	NULL	1000	750	0.00	0.00
3	3	BB Ball Bearing	BE-2349	1	0	NULL	800	600	0.00	0.00
4	4	Headset Ball Bearings	BE-2908	0	0	NULL	800	600	0.00	0.00
5	316	Blade	BL-2036	1	0	NULL	800	600	0.00	0.00
6	317	LL Crankarm	CA-5965	0	0	Black	500	375	0.00	0.00
7	318	ML Crankarm	CA-6738	0	0	Black	500	375	0.00	0.00
8	319	HL Crankarm	CA-7457	0	0	Black	500	375	0.00	0.00

4 + 100 + 50 + 1 + 1 + 30 + 2 + 2 + 8 + 8 = 201 bytes

Column widths affect physical buffer size...

201 bytes * 10,000 = 2MB < DefaultMaxBufferSize (10MB)
Estimated Row Size * DefaultMaxBufferRows

Buffers

Row 1

ProductID	Name	ProductNumber	MakeFlag	FinishedGoodsFlag	Color	SafetyStockLevel	ReorderPoint	StandardCost	ListPrice
x	xxxxxxx	xxxxxxx	x	x	xxxxx	xxx	xxx	xxx	xxxx
x	xxxxxxx	xxxxxxx	x	x	xxxxx	xxx	xxx	xxx	xxxx
x	xxxxxxx	xxxxxxx	x	x	xxxxx	xxx	xxx	xxx	xxxx
x	xxxxxxx	xxxxxxx	x	x	xxxxx	xxx	xxx	xxx	xxxx
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
50,000	x	xxxxxxx	x	x	xxxxx	xxx	xxx	xxx	xxxx

Row

50,001

ProductID	Name	ProductNumber	MakeFlag	FinishedGoodsFlag	Color	SafetyStockLevel	ReorderPoint	StandardCost	ListPrice
x	xxxxxxx	xxxxxxx	x	x	xxxxx	xxx	xxx	xxx	xxxx
x	xxxxxxx	xxxxxxx	x	x	xxxxx	xxx	xxx	xxx	xxxx
x	xxxxxxx	xxxxxxx	x	x	xxxxx	xxx	xxx	xxx	xxxx
x	xxxxxxx	xxxxxxx	x	x	xxxxx	xxx	xxx	xxx	xxxx
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
Row	x	xxxxxxx	x	x	xxxxx	xxx	xxx	xxx	xxxx
100,000									

...To a maximum of 5 buffers receiving data from a source

Non-Blocking Streaming Transformations

Audit

Character Map

Conditional Split

Copy Column

Data Conversion

Derived Column

Lookup (full cache)

Multicast

Percent Sampling

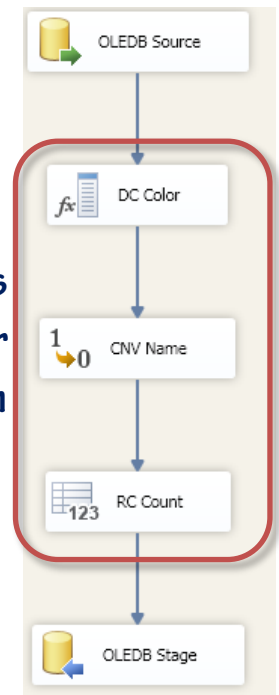
Row Count

Script Component

(without asynchronous output)

ProductID	Name	ProductNumber	MakeFlag	FinishedGoodsFlag
X	CCCCCCC	XXXXXXXX	X	X
X	CCCCCCC	XXXXXXXX	X	X
X	CCCCCCC	XXXXXXXX	X	X
X	CCCCCCC	XXXXXXXX	X	X
⋮	⋮	⋮	⋮	⋮
⋮	⋮	⋮	⋮	⋮
X	CCCCCCC	XXXXXXXX	X	X

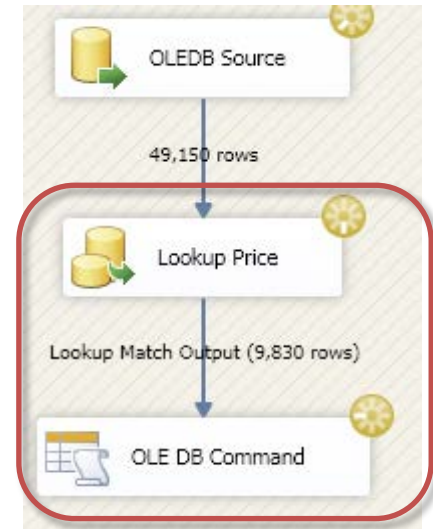
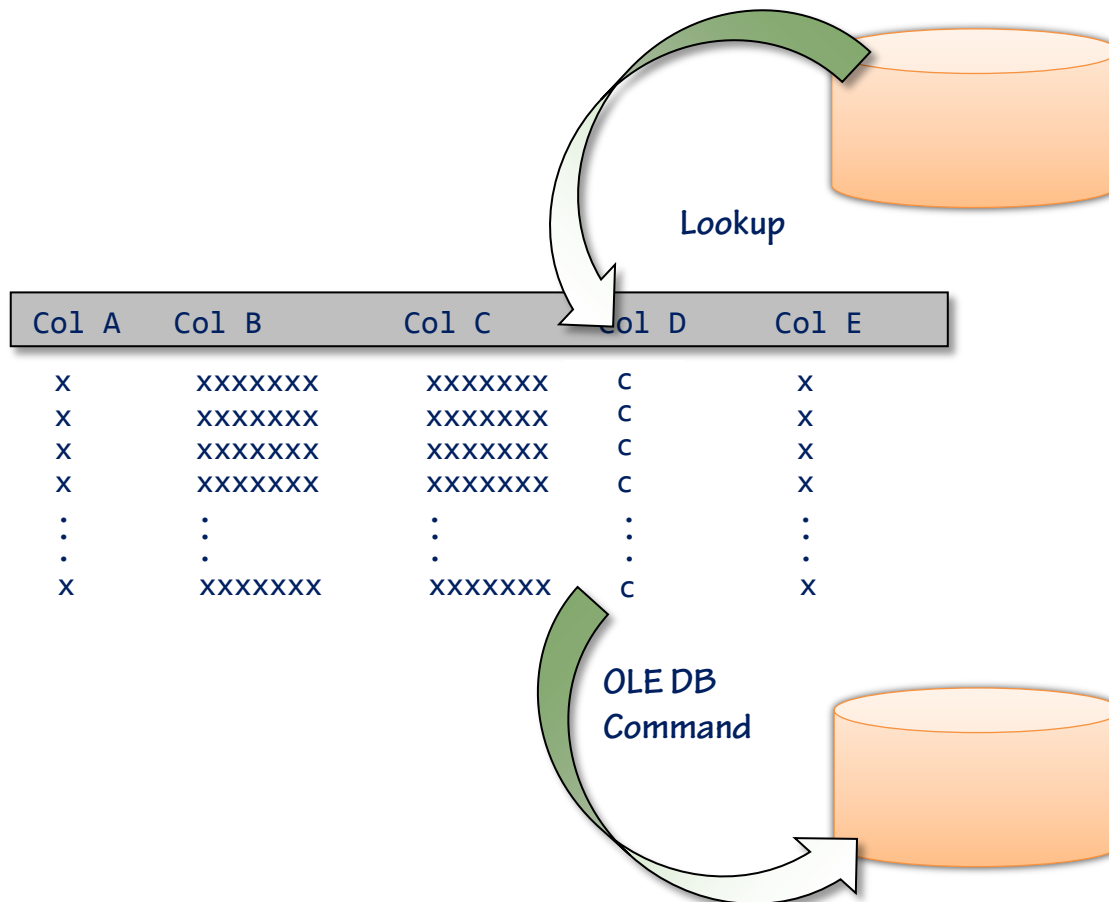
These transforms
operate on the same buffer
from start to finish



Non-Blocking Row-Based Transformations

Export Column
Import Column
Lookup (no cache or partial cache)

OLE DB Command
Script (with external interaction)
Slowly Changing Dimension



These transforms
operate row by row
but can't work on same
buffer simultaneously

Partially-Blocking Transformations

Data Mining Query

Merge

Merge Join

Pivot

Term Lookup

Unpivot

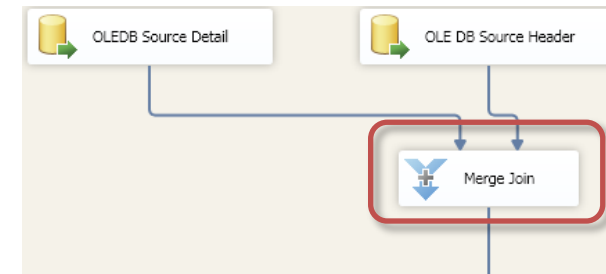
Union All

Col A	Col B	Col C	Col D
x	xxxxxxx	xxxxxxx	x
x	xxxxxxx	xxxxxxx	x
x	xxxxxxx	xxxxxxx	x
x	xxxxxxx	xxxxxxx	x
⋮	⋮	⋮	⋮
x	xxxxxxx	xxxxxxx	x

Col A	Col E
x	x
x	x
x	x
x	x
⋮	⋮
x	x

Merge Join

Col A	Col B	Col C	Col D	Col E
x	xxxxxxx	xxxxxxx	x	x
x	xxxxxxx	xxxxxxx	x	x
x	xxxxxxx	xxxxxxx	x	x
x	xxxxxxx	xxxxxxx	x	x
⋮	⋮	⋮	⋮	⋮
x	xxxxxxx	xxxxxxx	x	x



Partially-blocking
transforms
might wait on other
inputs and always
produce new buffer

Blocking Transformations

Aggregate

Fuzzy Grouping

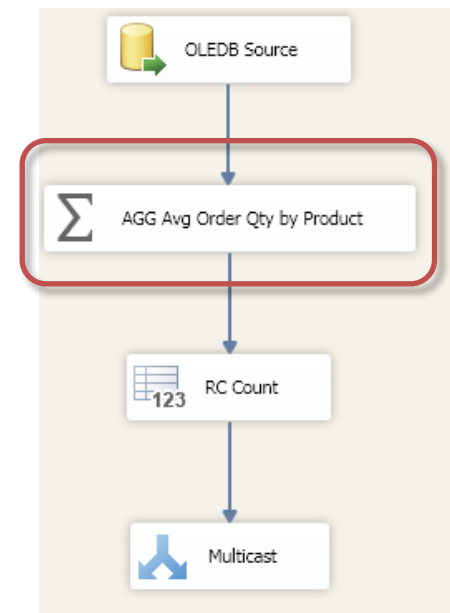
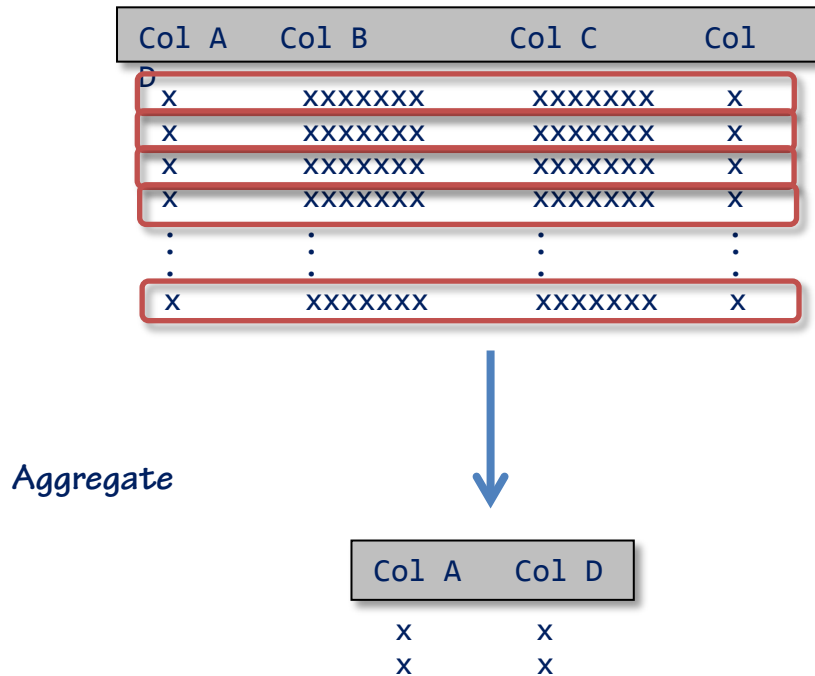
Fuzzy Lookup

Row Sampling

Sort

Term Extraction

Script Component (if code must read all rows)



These transforms stop pipeline until all rows read and always produce new buffer

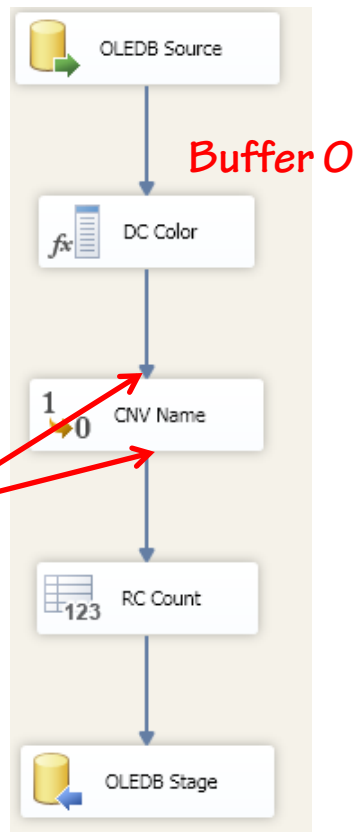
Synchronous versus Asynchronous Outputs

Non-blocking

Synchronous

Operate
on
Same buffer

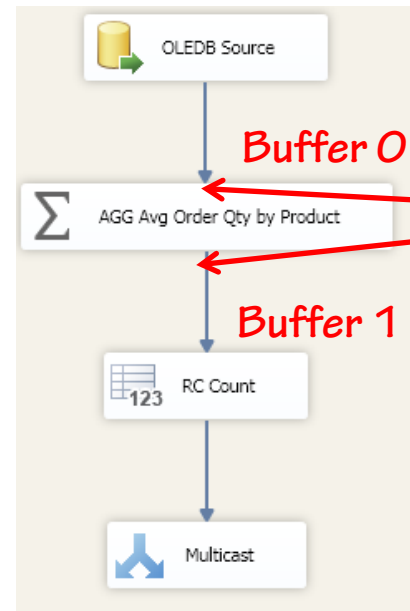
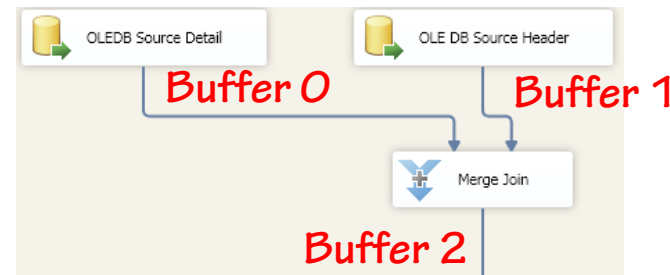
Input rows
=
Output rows



Partially blocking
Blocking

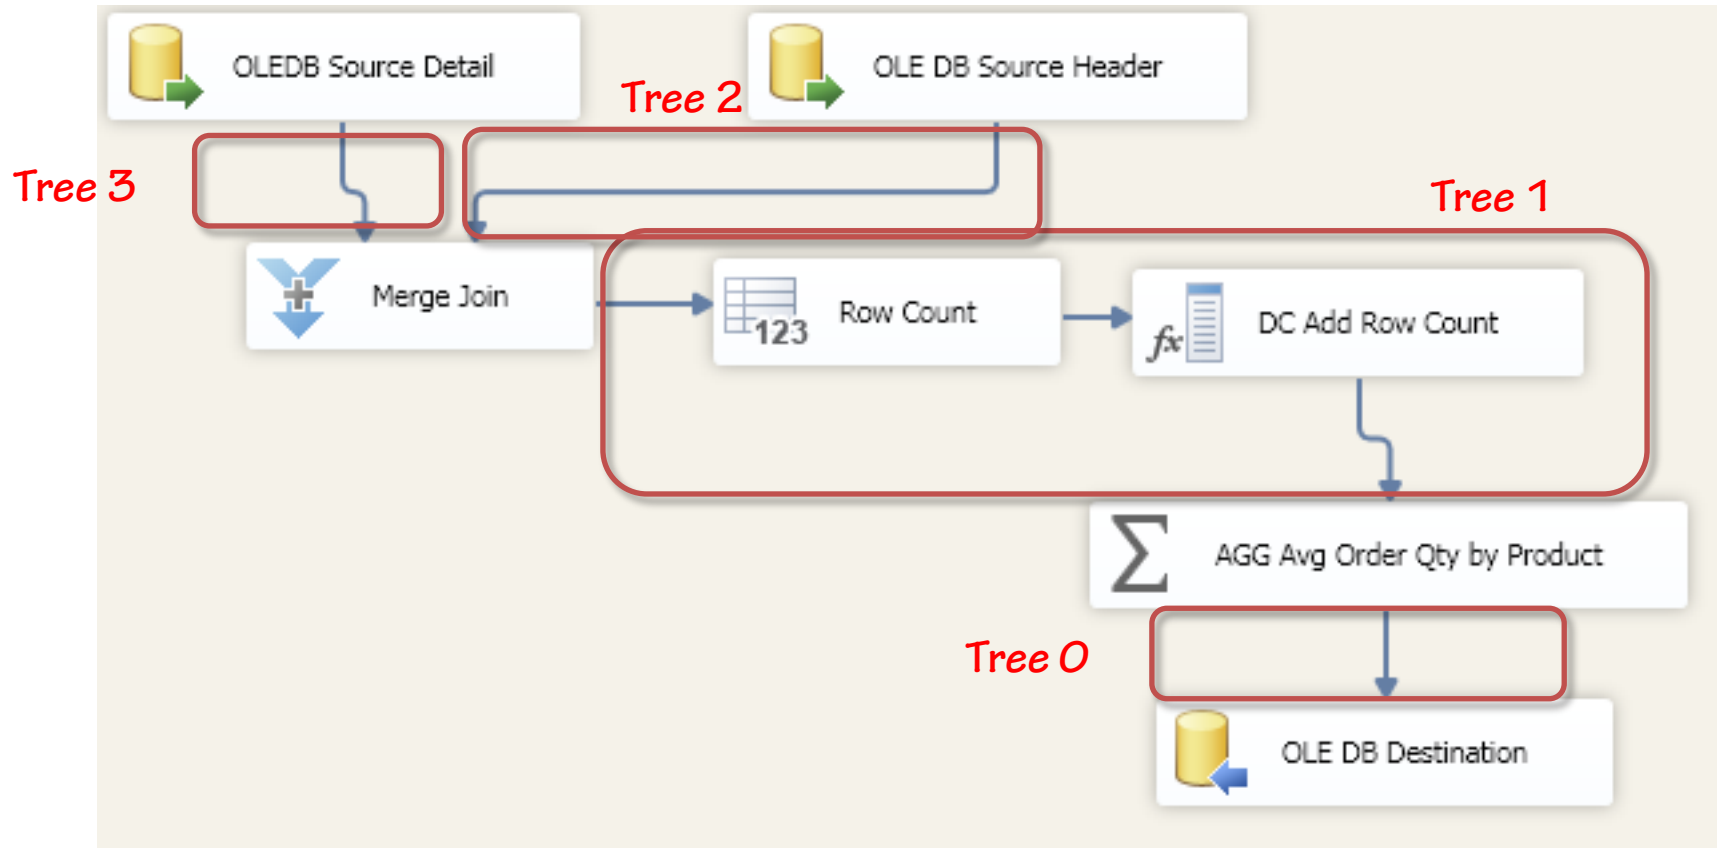
Asynchronous

Output into
New buffer

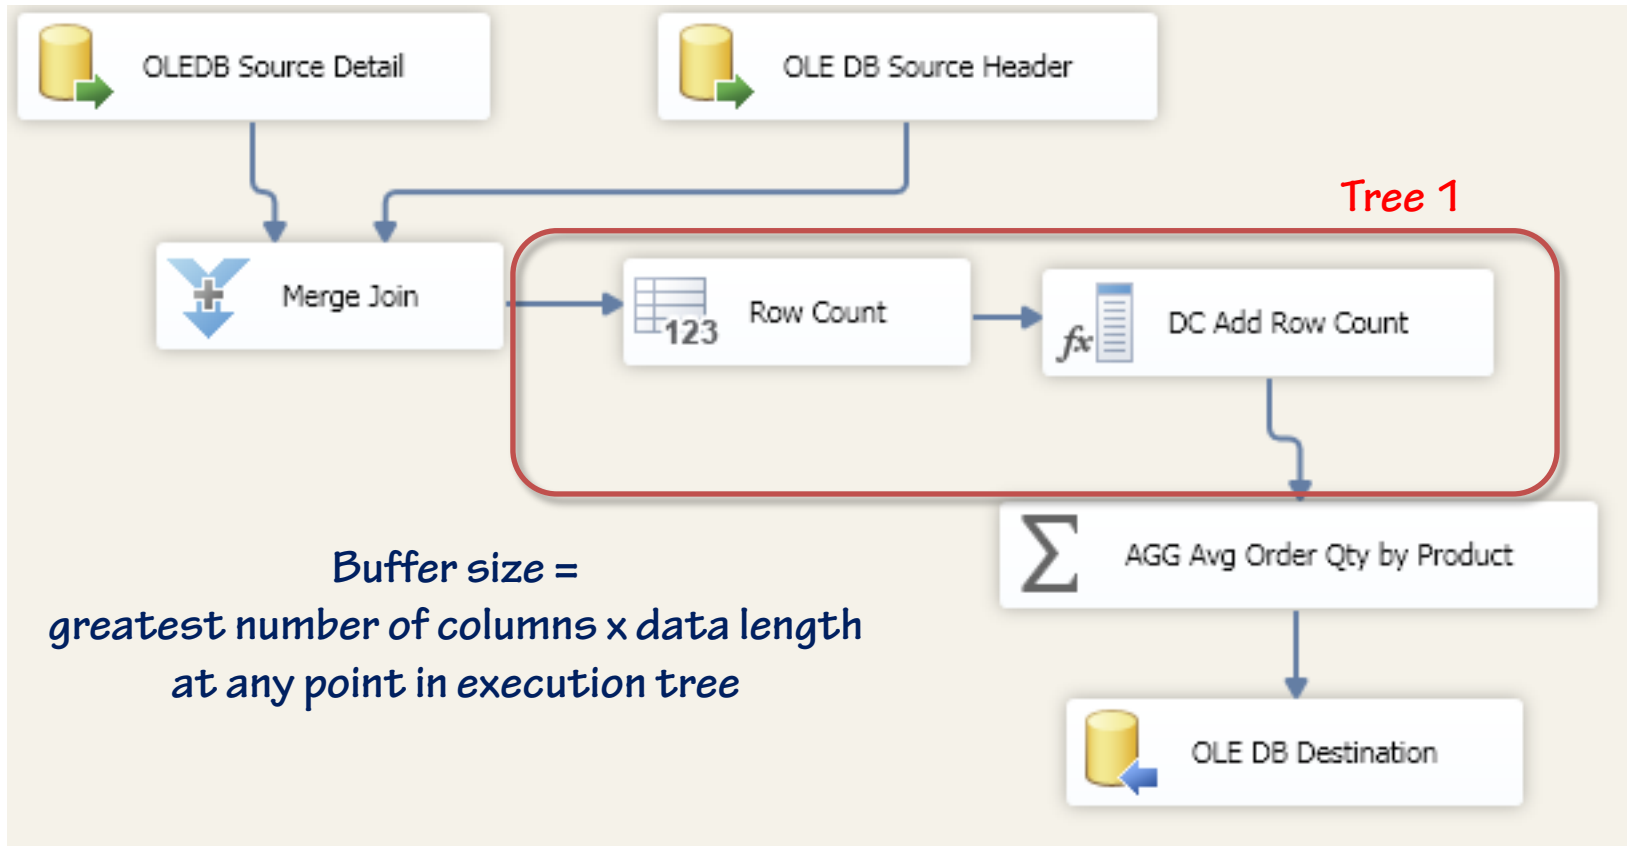


Input rows
<>
Output rows
(usually)

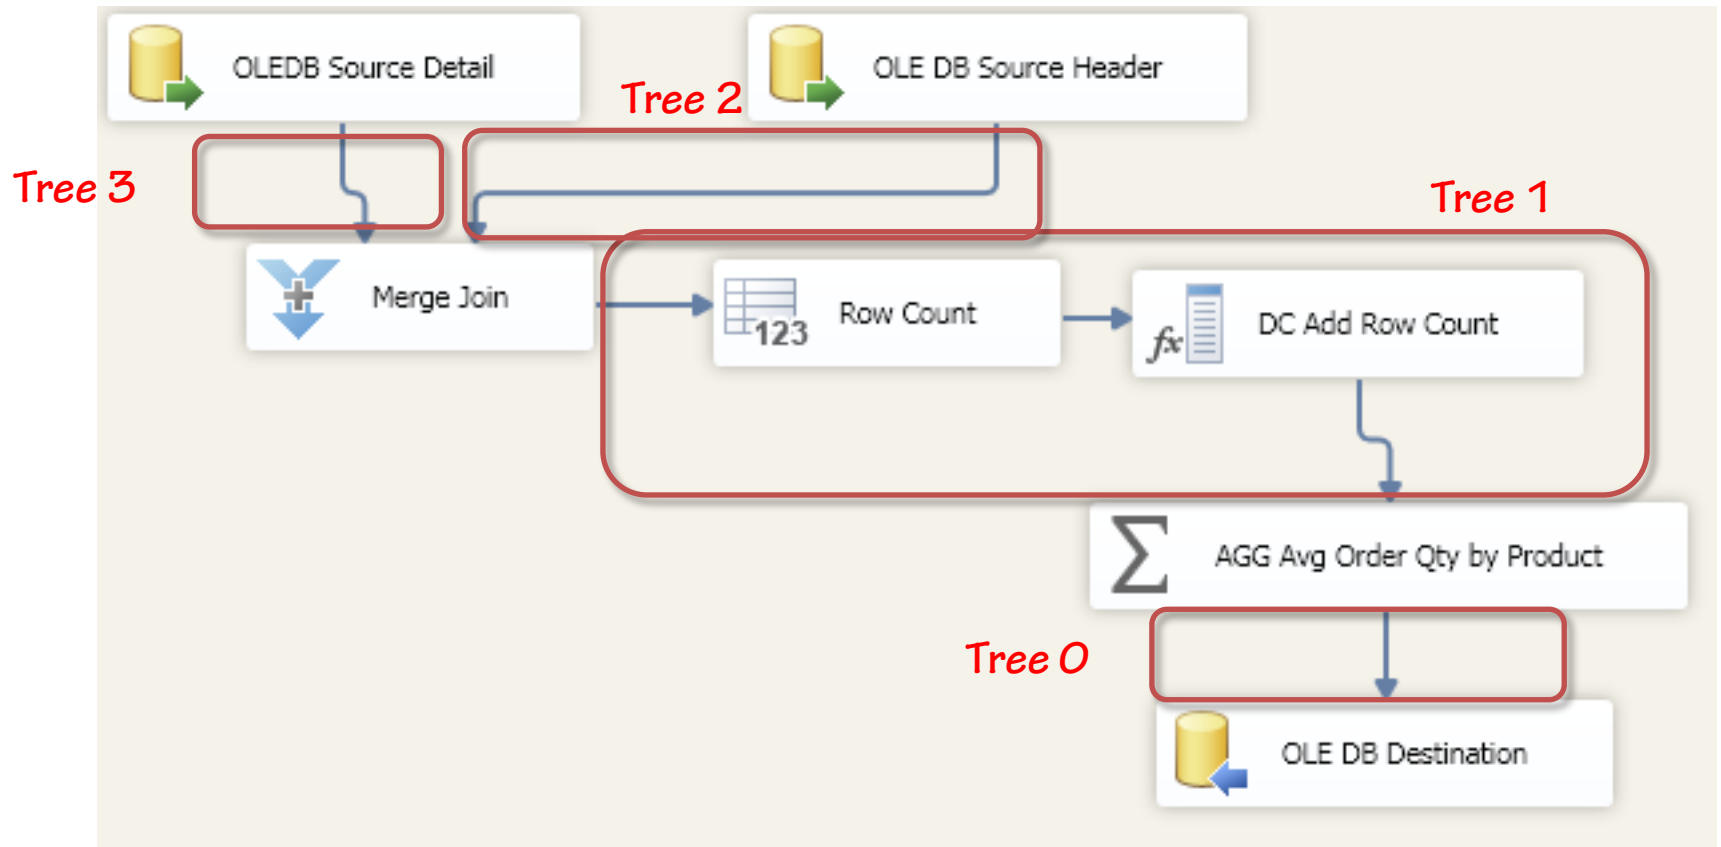
Execution Trees



Execution Trees



Execution Threads



Thread 0
Thread 1
Thread 2

Thread 3
Thread 4
Thread 5

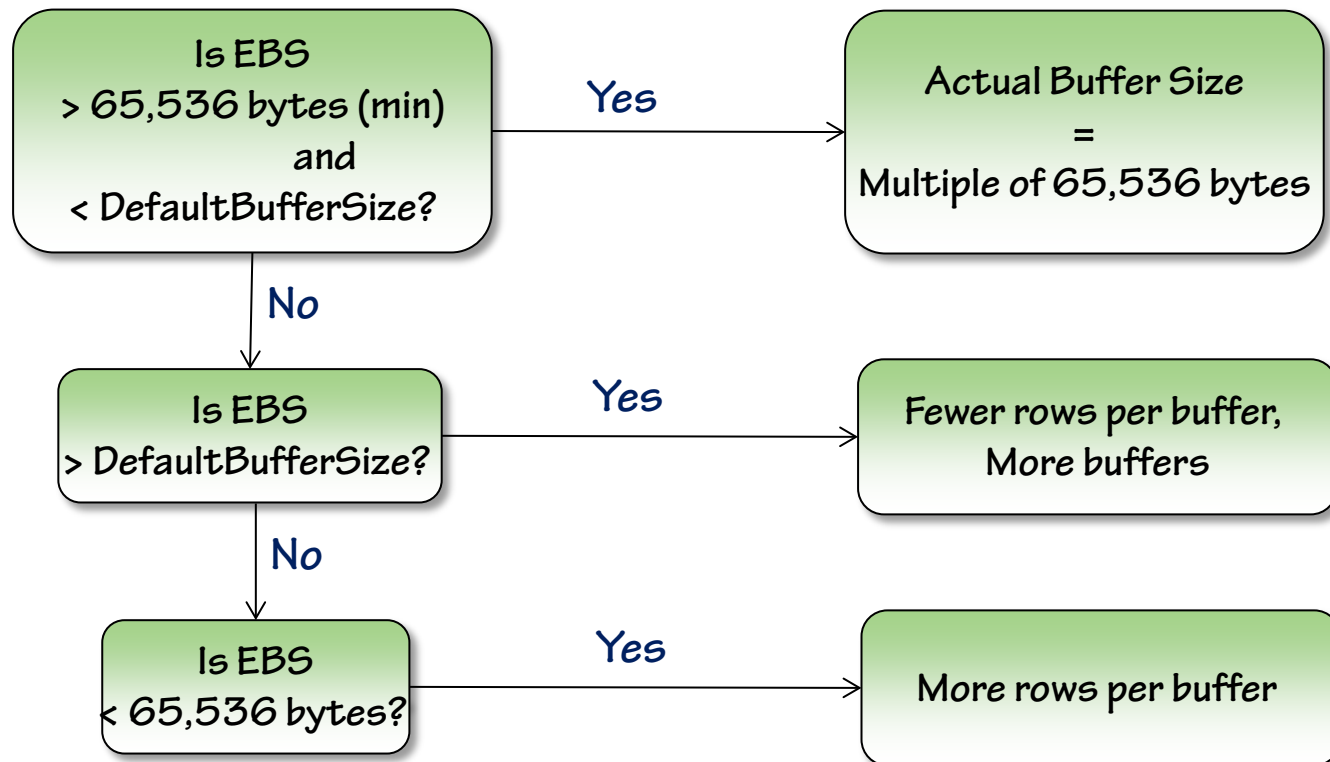
Increase parallelism by increasing EngineThreads value
(Data Flow Task property)
Minimum: # of Execution Trees + Sources

Performance Optimization: Buffer Sizing

Data Flow Task Properties

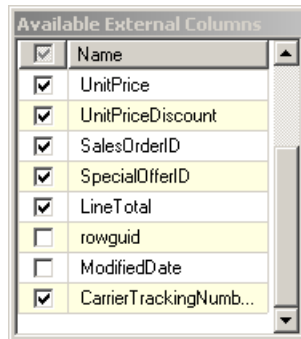
Misc	
BufferTempStoragePath	
DefaultBufferMaxRows	10000
DefaultBufferSize	10485760

$$\begin{array}{ccccc} \text{Estimated} & & & & \text{Estimated} \\ \text{Row} & & \text{Max} & & \text{Buffer} \\ \text{Size} & \times & \text{Buffer} & = & \text{Size} \\ & & \text{Rows} & & \text{(EBS)} \end{array}$$




Performance Optimization: Buffer Sizing


- Reduce the number of columns and use smallest data types possible





*Use query conversions to reduce data length
if possible*


- Warnings alert you to unused columns in data flow

 [SSIS.Pipeline] Warning: The output column "SalesOrderDetailID" (17) on output "OLE DB Source Output" (13) and component "OLEDB Source" (2) is not subsequently used in the Data Flow task.

 [SSIS.Pipeline] Warning: The output column "CarrierTrackingNumber" (18) on output "OLE DB Source Output" (13) and component "OLEDB Source" (2) is not subsequently used in the Data Flow task.

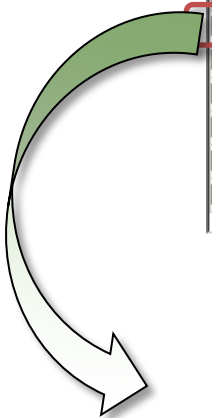
 [SSIS.Pipeline] Warning: The output column "OrderQty" (19) on output "OLE DB Source Output" (13) and component "OLEDB Source" (2) is not subsequently used in the Data Flow task.

 [SSIS.Pipeline] Warning: The output column "ProductID" (20) on output "OLE DB Source Output" (13) and component "OLEDB Source" (2) is not subsequently used in the Data Flow task.

 [SSIS.Pipeline] Warning: The output column "SpecialOfferID" (21) on output "OLE DB Source Output" (13) and component "OLEDB Source" (2) is not subsequently used in the Data Flow task.

Performance Optimization: Buffer Sizing

- Log package with BufferSizeTuning to see adjustment of buffers



<input checked="" type="checkbox"/>	Events	Description
<input type="checkbox"/>	OnPostExecute	Handles post-execution events. Use to define post-...
<input type="checkbox"/>	OnPostValidate	Handles post-validation events. Use to define post-...
<input type="checkbox"/>	OnPreExecute	Handles pre-execution events. Use to define pre-pr...
<input type="checkbox"/>	OnPreValidate	Handles pre-validation events. Use to define pre-pr...
<input type="checkbox"/>	OnProgress	Handles progress notifications. Use to define action...
<input type="checkbox"/>	OnQueryCancel	Handles cancel events. Called periodically to determ...
<input type="checkbox"/>	OnTaskFailed	Handles task failures. Use to define actions to perfo...
<input type="checkbox"/>	OnVariableValueChanged	Handles value changes in variables whose RaiseCha...
<input type="checkbox"/>	OnWarning	Handles warning events. Use to define actions whe...
<input checked="" type="checkbox"/>	BufferSizeTuning	The Data Flow changed the size of a buffer away fr...
<input type="checkbox"/>	PipelineBufferLeak	After execution, some buffers were left outstandin...
<input type="checkbox"/>	PipelineComponentTime	Information about each phase of validation and exe...
<input type="checkbox"/>	PipelineExecutionPlan	The execution plan of the data flow.
<input type="checkbox"/>	PipelineExecutionTrees	The input the scheduler had when forming the exec...
<input type="checkbox"/>	PipelineInitialization	Various information from Data Flow initialization.

Log the data flow task

User:BufferSizeTuning DFT Transform-Load
Rows in buffer type 5 would cause a buffer size greater than the configured maximum.

There will be only 15065 rows in buffers of this type.

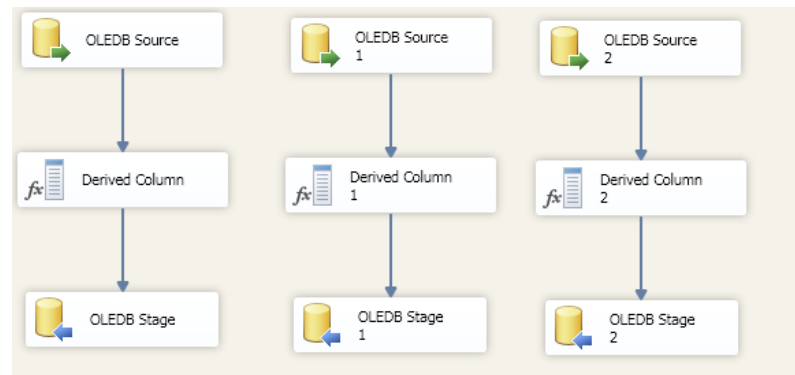
Package Performance: Parallelism

- Increase MaxConcurrentExecutables on dedicated machine

Execution	
DelayValidation	False
Disable	False
DisableEventHandlers	False
FailPackageOnFailure	False
FailParentOnFailure	False
MaxConcurrentExecutables	-1
MaximumErrorCount	1
PackagePriorityClass	Default

-1 = Number of processors + 2

- Use parallel pipelines in data flow with separate data sets



Package Performance: Transformations


- Choose transformations carefully

Best choice: Non-blocking


Audit
Character Map
Conditional Split
Copy Column
Data Conversion
Derived Column

Lookup (full cache)
Multicast
Percent Sampling
Row Count
Script Component
(without asynchronous output)

- Reduce rows with Conditional Split where possible

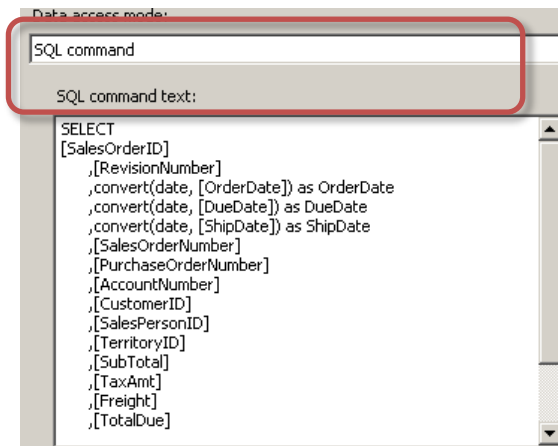


- Optimize the Lookup cache



Package Performance: Sources

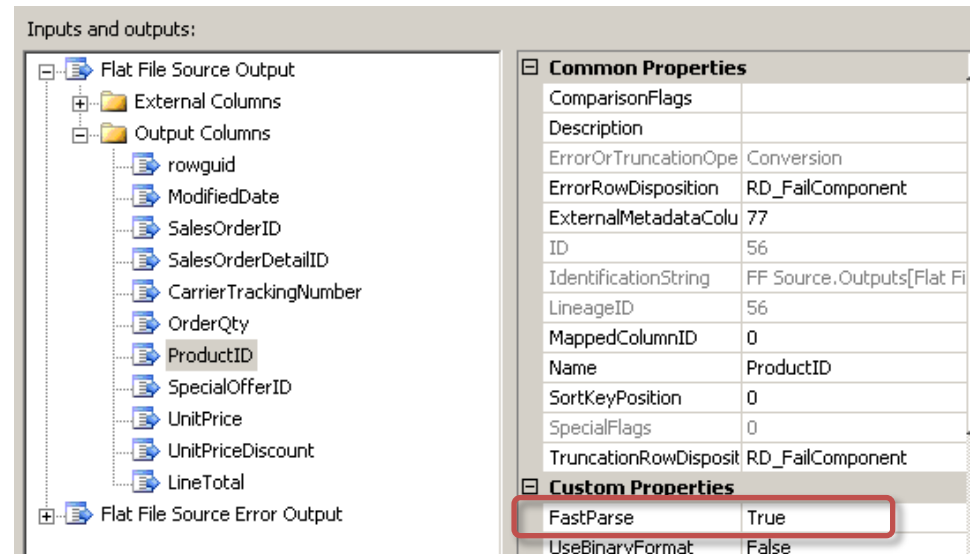
- Data Access Mode: SQL Command for relational sources



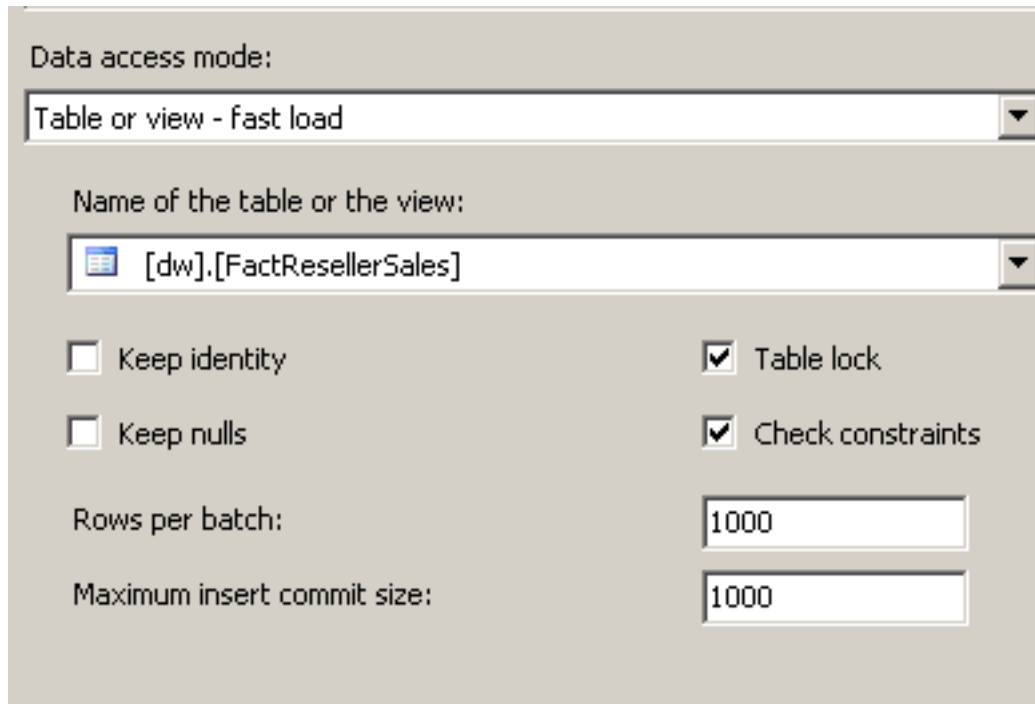
*Uses sp_executesql
(faster than OPENROWSET)*

- FastParse for flat file sources

*Use for when data types
Not required to be locale-specific*



Package Performance: OLE DB Destinations



The screenshot shows the 'Data access mode' dialog box. The 'Data access mode' dropdown is set to 'Table or view - fast load'. The 'Name of the table or the view' dropdown is set to '[dw].[FactResellerSales]'. The 'Keep identity' checkbox is unchecked, while 'Keep nulls', 'Table lock', and 'Check constraints' are checked. The 'Rows per batch' and 'Maximum insert commit size' are both set to 1000.

Property	Value
Data access mode	Table or view - fast load
Name of the table or the view	[dw].[FactResellerSales]
Keep identity	<input type="checkbox"/>
Keep nulls	<input type="checkbox"/>
Table lock	<input checked="" type="checkbox"/>
Check constraints	<input checked="" type="checkbox"/>
Rows per batch	1000
Maximum insert commit size	1000

Bulk insert with Fast Load option

Table lock performs faster

Clustered index?

Set rows per batch for faster inserts

Or...

Drop the index, load rows, and rebuild index

Performance Monitoring: Memory

Memory

Page Reads/sec

0.000

Page Reads/sec > 500
= memory pressure

Network Interface

Bytes Total/sec

Intel[R] PRO_1000 MT Network Connection _2

0.000

Process

% Processor Time

MsDtsSrvr

0.000

sqlservr

0.000

Private Bytes

51,306,496.000

Working Set

22,401,024.000

SQLServer:Memory Manager

Total Server Memory (KB)

1,264,768.000

SQLServer:SSIS Pipeline 11.0

Buffer memory

13,835,752.000

Buffers in use

17.000

Buffers spooled

0.000

Flat buffer memory

8,232.000

Flat buffers in use

3.000

Performance Monitoring: CPU

Memory

Page Reads/sec 0.000

Network Interface

Intel[R] PRO_1000 MT Network Connection _2

Bytes Total/sec 0.000

Process	MsDtsSrvr	sqlservr
% Processor Time	0.000	0.000
Private Bytes	51,306,496.000	---
Working Set	22,401,024.000	---

SQLServer:Memory Manager

Total Server Memory (KB) 1,264,768.000

SQLServer:SSIS Pipeline 11.0

Buffer memory 13,835,752.000

Buffers in use 17.000

Buffers spooled 0.000

Flat buffer memory 8,232.000

Flat buffers in use 3.000

Performance Monitoring: Network

Memory

Page Reads/sec	0.000
----------------	-------

Network Interface	Intel[R] PRO_1000 MT Network Connection _2
Bytes Total/sec	0.000

Process	MsDtsSrvr	sqlservr
% Processor Time	0.000	0.000
Private Bytes	51,306,496.000	---
Working Set	22,401,024.000	---

SQLServer:Memory Manager	
Total Server Memory (KB)	1,264,768.000

SQLServer:SSIS Pipeline 11.0	
Buffer memory	13,835,752.000
Buffers in use	17.000
Buffers spooled	0.000
Flat buffer memory	8,232.000
Flat buffers in use	3.000

Performance Monitoring: Integration Services

Memory

Page Reads/sec	0.000
----------------	-------

Network Interface

Bytes Total/sec	Intel[R] PRO_1000 MT Network Connection _2	0.000
-----------------	--	-------

Process

	MsDtsSrvr	sqlservr
% Processor Time	0.000	0.000
Private Bytes	51,306,496.000	---
Working Set	22,401,024.000	---

SQLServer:Memory Manager

Total Server Memory (KB)	1,264,768.000
--------------------------	---------------

SQLServer:SSIS Pipeline 11.0

Buffer memory	13,835,752.000
Buffers in use	17.000
Buffers spooled	0.000
Flat buffer memory	8,232.000
Flat buffers in use	3.000

Summary

- **Pipeline Buffers**

- Buffers, transformations, synchronous and asynchronous outputs

- **Execution Trees and Execution Threads**

- Asynchronous output starts new execution tree and allocation of at least one thread

- **Package Performance Optimization**

- Buffer size, parallelism, transformation selection, and adapter optimization (sources, destinations)

- **Performance Monitoring**

- Memory, CPU, network, Integration Services buffers (especially buffer spool)

Resources

- **Top 10 SQL Server Integration Services Best Practices**
 - *<http://tinyurl.com/7hasrbm>*
- **Integration Services: Performance Tuning Techniques**
 - *<http://tinyurl.com/lzaeszr>*
- **Integration Services Performance Counters**
 - *<http://tinyurl.com/lk57stq>*